

Congratulations! You passed!

TO PASS 80% or higher

Keep Learning

GRADE

100%

Natural Language Processing & Word Embeddings

LATEST SUBMISSION GRADE

100%

1.

Question 1

Suppose you learn a word embedding for a vocabulary of 10000 words. Then the embedding vectors should be 10000 dimensional, so as to capture the full range of variation and meaning in those words.

1 / 1 point



True



False

Correct

The dimension of word vectors is usually smaller than the size of the vocabulary. Most common sizes for word vectors range between 50 and 400.

2.

Question 2

What is t-SNE?

1 / 1 point



A non-linear dimensionality reduction technique



A linear transformation that allows us to solve analogies on word vectors



A supervised learning algorithm for learning word embeddings



An open-source sequence modeling library

Correct

Yes

3.

Question 3

Suppose you download a pre-trained word embedding which has been trained on a huge corpus of text. You then use this word embedding to train an RNN for a language task of recognizing if someone is happy from a short snippet of text, using a small training set.

x (input text)	y (happy?)
I'm feeling wonderful today!	1
I'm bummed my cat is ill.	0
Really enjoying this!	1

Then even if the word “ecstatic” does not appear in your small training set, your RNN might reasonably be expected to recognize “I’m ecstatic” as deserving a label $y = 1$.

1 / 1 point



True



False

Correct

Yes, word vectors empower your model with an incredible ability to generalize. The vector for “ecstatic” would contain a positive/happy connotation which will probably make your model classify the sentence as a “1”.

4.

Question 4

Which of these equations do you think should hold for a good word embedding? (Check all that apply)

1 / 1 point



$e_{\text{boy}} - e_{\text{girl}} \approx e_{\text{sister}} - e_{\text{brother}}$ $e_{\text{boy}} - e_{\text{girl}} \approx e_{\text{sister}} - e_{\text{brother}}$



$e_{\text{boy}} - e_{\text{girl}} \approx e_{\text{brother}} - e_{\text{sister}}$ $e_{\text{boy}} - e_{\text{girl}} \approx e_{\text{brother}} - e_{\text{sister}}$

Correct

Yes!



$e_{\text{boy}} - e_{\text{brother}} \approx e_{\text{girl}} - e_{\text{sister}}$ $e_{\text{boy}} - e_{\text{brother}} \approx e_{\text{girl}} - e_{\text{sister}}$

Correct

Yes!



$e_{\text{boy}} - e_{\text{brother}} \approx e_{\text{sister}} - e_{\text{girl}}$ $e_{\text{boy}} - e_{\text{brother}} \approx e_{\text{sister}} - e_{\text{girl}}$

5.

Question 5

Let E be an embedding matrix, and let o_{1234} be a one-hot vector corresponding to word 1234. Then to get the embedding of word 1234, why don't we call $E * o_{1234}$ in Python?

1 / 1 point



The correct formula is $E^T * o_{1234}$.



None of the above: calling the Python snippet as described above is fine.



This doesn't handle unknown words (<UNK>).



It is computationally wasteful.

Correct

Yes, the element-wise multiplication will be extremely inefficient.

6.

Question 6

When learning word embeddings, we create an artificial task of estimating $P(\text{target} \mid \text{context})$. It is okay if we do poorly on this artificial prediction task; the more important by-product of this task is that we learn a useful set of word embeddings.

1 / 1 point



False



True

Correct

7.

Question 7

In the word2vec algorithm, you estimate $P(t \mid c)P(t|c)$, where tt is the target word and cc is a context word. How are tt and cc chosen from the training set? Pick the best answer.

1 / 1 point



cc is a sequence of several words immediately before tt .



cc is the sequence of all the words in the sentence before tt .



cc is the one word that comes immediately before tt .



cc and tt are chosen to be nearby words.

Correct

8.

Question 8

Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The word2vec model uses the following softmax function:

$$P(t \mid c) = \frac{e^{\theta_t^T e_c}}{\sum_{t'=1}^{10000} e^{\theta_{t'}^T e_c}} P(t|c) = \sum_{t'=1}^{10000} e^{\theta_{t'}^T e_c} e^{\theta_t^T e_c}$$

Which of these statements are correct? Check all that apply.

1 / 1 point



θ_t and e_c are both 10000 dimensional vectors.



θ_t and e_c are both trained with an optimization algorithm such as Adam or gradient descent.

Correct



θ_t and e_c are both 500 dimensional vectors.

Correct



After training, we should expect θ_t to be very close to e_c when t and c are the same word.

9.

Question 9

Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The GloVe model minimizes this objective:

$$\min \sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(X_{ij}) (\theta_i^T e_j + b_i + b_j - \log X_{ij})^2$$

Which of these statements are correct? Check all that apply.

1 / 1 point



The weighting function $f(\cdot)$ must satisfy $f(0) = 0$.

Correct

The weighting function helps prevent learning only from extremely common word pairs. It is not necessary that it satisfies this function.



X_{ij} is the number of times word j appears in the context of word i .

Correct



θ_i and e_j should be initialized to 0 at the beginning of training.



θ_i and e_j should be initialized randomly at the beginning of training.

Correct

10.

Question 10

You have trained word embeddings using a text dataset of m_1 words. You are considering using these word embeddings for a language task, for which you have a separate labeled dataset of m_2 words. Keeping in mind that using word embeddings is a form of transfer

learning, under which of these circumstances would you expect the word embeddings to be helpful?

1 / 1 point



$m_1 \ll m_2$



$m_1 \gg m_2$

Correct